

Understanding The Problem Statement(18):

Who deals with this:

Answering RFI's usually involves people from different categories ,each doing their own research without knowing what others have already figured out.

The actual problem and our way to solve it:

There's no central place where past RFI's and their approved answers are saved in an organized way. So whenever a similar issue comes up ,teams end up solving the same issue from scratch and we are building a shared repository for RFI's ,their considerations and responses .

Why it is concerning:

The real issue is that nothing is saved in a way that's easy to find and reuse. And because this happens across roughly 650 applications a year, even small, avoidable delays add up to a significant, recurring cost in time and effort.

FEATURE 1(HYBRID SEARCH):

How it's stored:

Every RFI is stored in three layers. The original document as a file, the structured details like section, category, and date as database fields, and the meaning of the text as a vector embedding. This means you can find past RFIs by browsing, filtering, or searching by meaning — all three ways work.

What it does:

Type in an RFI, and the system runs semantic search(matching by meaning), not just exact words combined with keyword search for exact terms like regulation codes. Together, that's hybrid search, so nothing gets missed either way.

What you get back:

The closest matching past RFIs, their approved responses, and a confidence score showing how closely each result matches your query.

FEATURE 2(PROACTIVE RISK SCORING):

What it does:

Upload a draft application before submitting it. The system checks each section against patterns from past sections that triggered RFIs and flags the risky ones, before a regulator ever sees them.

What you get back:

A risk score for each section, not just one score for the whole thing. For example Protocol Design comes back Low, IMPD Quality comes back High. Along with each score, it shows the past RFI it resembles and the response that worked for it, so the team can fix it in advance.

Why it matters:

This is the difference between reacting to an RFI and preventing one. Every other tool helps you answer faster after the fact. This catches the problem before it's ever sent.

FEATURE 3(DRAFT RESPONSE GENERATION):

Once an RFI arrives the system uses a RAG pipeline, It first retrieves the most similar past RFIs and approved responses from the repository, then passes them to Gemini to generate a ready to review draft grounded in real precedent. The AI does not guess. It references. Differences between the current and past RFIs are flagged so reviewers know exactly what to double check before approving.

Feature 4 (Analytics Dashboard):

Tracks RFI patterns across all 650 submissions identifying which sections trigger the most RFIs and where submission quality can be permanently improved. Turns individual fixes into organisation wide learning.

Feature 5 (Audit Trail):

Every response goes through a tracked draft, review, approve, and submit workflow. The audit trail is divided by section and by team. Protocol Design, IMPD Quality, and Clinical Overview each have their own history. RA Clinical, Affiliates, CTA Management, and the EU Submission Hub each see only their own work. Each approved response automatically feeds back into the repository making the system smarter over time.